

Reinforcement Learning for Optimal Trade Scheduling using Soft Actor-Critic (SAC) Model

Prem Shenoy

October 26, 2024

1 Introduction

In algorithmic trading, executing large orders while minimizing transaction costs is essential. Traditional methods, such as *Time-Weighted Average Price* (TWAP) and *Volume-Weighted Average Price* (VWAP), have limitations as they are static and do not adapt to market conditions. In contrast, reinforcement learning models, such as **Soft Actor-Critic (SAC)**, enable dynamic decision-making, optimizing the order execution process in real-time. This report presents an analysis of SAC-based trade scheduling and compares its performance against TWAP and VWAP strategies.

2 Soft Actor-Critic (SAC) Model

SAC is a state-of-the-art, model-free reinforcement learning algorithm used for continuous control tasks. It maximizes expected reward while optimizing entropy, allowing it to select more stochastic actions. This is beneficial in dynamic environments like financial markets, where conditions fluctuate continuously.

In this problem, SAC dynamically determines the optimal number of shares to sell at each time step while minimizing the costs associated with:

- **Slippage:** The difference between the expected execution price and the actual execution price.
- **Market Impact:** The effect of large orders on the price of the asset.

The reinforcement learning environment is designed with a trading horizon of 390 minutes (one trading day), and the goal is to sell 1000 shares of AAPL.

3 Experiment Setup

We simulate the market environment using historical AAPL data. The SAC agent is trained to determine the number of shares to sell at each minute, while TWAP and VWAP strategies are benchmarked for comparison.

3.1 TWAP and VWAP Benchmarks

- **TWAP:** This strategy sells an equal number of shares at every time step throughout the trading period.
- **VWAP:** This strategy sells shares in proportion to the trading volume at each time step.

3.2 Metrics for Comparison

We compare the strategies based on their total transaction costs, which include both slippage and market impact.

4 Results and Analysis

4.1 Training Dynamics of SAC

The SAC model was trained over 264 episodes, accumulating a total of 99,724 time steps. During training, various statistics were collected, including:

- **Actor Loss:** The loss incurred by the policy network, which was 7.5×10^3 .
- **Critic Loss:** The loss of the value function, increasing with the complexity of the learned policy, with a final value of 1.91×10^8 .
- **Entropy Coefficient:** A parameter that governs exploration vs. exploitation, starting at 1.03 and reaching -0.721 .

4.2 Cost Comparison

After training the SAC model, we compared the total costs of executing the trade schedule against the TWAP and VWAP strategies. The results are summarized below:

- **SAC Total Cost:** \$170,293.81
- **TWAP Total Cost:** \$184,357.87
- **VWAP Total Cost:** \$5,020.39

4.2.1 SAC Analysis

Remarkably, the SAC model managed to reduce the total transaction cost by 7.63% when compared to TWAP, but is unable to beat VWAP. This indicates that the dynamic trade scheduling learned by the SAC model was able to avoid some of the adverse effects of both slippage and market impact which can be attributed to the model's ability to dynamically adjust its strategy based on real-time market data, but additional fine-tuning of the hyper-parameters is needed to further reduce costs.

4.2.2 TWAP Analysis

The TWAP strategy incurred a significantly higher transaction cost of \$184,357.87. This high cost stems from TWAP's rigid structure, which does not adapt to market conditions. As a result, it led to unfavorable trade executions, particularly during periods of high volatility or low liquidity.

4.2.3 VWAP Analysis

VWAP, being volume-adaptive, outperformed both TWAP and SAC, with a total cost of \$5,020.39.

5 Conclusion

This report demonstrates the effectiveness of using reinforcement learning, specifically the Soft Actor-Critic model, for optimizing trade scheduling. The SAC model outperformed traditional TWAP strategy, but was unable to beat VWAP. It was able to reduce the baseline cost by 7.63% compared to TWAP, but additional fine-tuning is needed to beat the VWAP baseline. In future work, we could explore the integration of other reinforcement learning models, such as Proximal Policy Optimization (PPO), to further improve the robustness of trade scheduling in highly volatile markets.

6 Appendix: Code

```
import gym
import numpy as np
import pandas as pd
from stable_baselines3 import SAC, PPO
from stable_baselines3.common.vec_env import DummyVecEnv
from gym import spaces
```

```

class Benchmark:
    def __init__(self, data):
        """
        Initializes the Benchmark class with provided market data.
        """
        self.data = data

    def get_twap_trades(self, data, initial_inventory, preferred_timeframe=390):
        total_steps = len(data)
        twap_shares_per_step = initial_inventory / preferred_timeframe
        remaining_inventory = initial_inventory
        trades = []
        for step in range(min(total_steps, preferred_timeframe)):
            size_of_slice = min(twap_shares_per_step, remaining_inventory)
            remaining_inventory -= int(np.ceil(size_of_slice))
            trade = {
                'timestamp': data.iloc[step]['timestamp'],
                'step': step,
                'price': data.iloc[step]['bid_price_1'],
                'shares': size_of_slice,
                'inventory': remaining_inventory,
            }
            trades.append(trade)
        return pd.DataFrame(trades)

    def get_vwap_trades(self, data, initial_inventory, preferred_timeframe=390):
        total_volume = data[['bid_size_1', 'bid_size_2', 'bid_size_3', 'bid_size_4', 'bid_size_5']].sum()
        total_steps = len(data)
        remaining_inventory = initial_inventory
        trades = []
        for step in range(min(total_steps, preferred_timeframe)):
            volume_at_step = data.iloc[step]['bid_size_1'] + data.iloc[step]['bid_size_2'] + data.iloc[step]['bid_size_3'] + data.iloc[step]['bid_size_4'] + data.iloc[step]['bid_size_5']
            size_of_slice = (volume_at_step / total_volume) * initial_inventory
            size_of_slice = min(size_of_slice, remaining_inventory)
            remaining_inventory -= int(np.ceil(size_of_slice))
            trade = {
                'timestamp': data.iloc[step]['timestamp'],
                'step': step,
                'price': data.iloc[step]['bid_price_1'],
                'shares': size_of_slice,
                'inventory': remaining_inventory,
            }
            trades.append(trade)
        return pd.DataFrame(trades)

    def calculate_vwap(self, idx, shares):
        bid_prices = self.data.iloc[idx][['bid_price_1', 'bid_price_2', 'bid_price_3', 'bid_price_4', 'bid_price_5']]
        bid_sizes = self.data.iloc[idx][['bid_size_1', 'bid_size_2', 'bid_size_3', 'bid_size_4', 'bid_size_5']]
        cumsum = 0
        for i, size in enumerate(bid_sizes):
            cumsum += size
            if cumsum >= shares:
                break
        return np.sum(bid_prices[:i+1] * bid_sizes[:i+1]) / np.sum(bid_sizes[:i+1])

    def compute_components(self, alpha, shares, idx):
        actual_price = self.calculate_vwap(idx, shares)

```

```

        Slippage = (self.data.iloc[idx]['bid_price_1'] - actual_price) * shares
        Market_Impact = alpha * np.sqrt(shares)
        return np.array([Slippage, Market_Impact])

def simulate_strategy(self, trades, data, preferred_timeframe):
    slippage = []
    market_impact = []
    alpha = 4.439584265535017e-06

    for idx in range(len(trades)):
        shares = trades.iloc[idx]['shares']
        reward = self.compute_components(alpha, shares, idx)
        slippage.append(reward[0])
        market_impact.append(reward[1])

    return slippage, market_impact

class TradingEnv(gym.Env):
    def __init__(self, data, total_shares=1000, trading_horizon=390, benchmark=None):
        super(TradingEnv, self).__init__()

        self.data = data
        self.total_shares = total_shares
        self.trading_horizon = trading_horizon
        self.current_step = 0
        self.remaining_shares = total_shares
        self.benchmark = benchmark

        self.action_space = spaces.Box(low=0, high=1, shape=(1,), dtype=np.float32)
        self.observation_space = spaces.Box(low=0, high=np.inf, shape=(2,), dtype=np.float32)

    def reset(self):
        self.current_step = 0
        self.remaining_shares = self.total_shares
        return self._get_observation()

    def _get_observation(self):
        current_price = self.data.iloc[self.current_step]['bid_price_1']
        return np.array([current_price, self.remaining_shares], dtype=np.float32)

    def step(self, action):
        current_price = self.data.iloc[self.current_step]['bid_price_1']
        current_volume = self.data.iloc[self.current_step][['bid_size_1', 'bid_size_2', 'bid_size_3']]

        shares_to_sell = max(0.01, action[0]) * self.remaining_shares
        shares_to_sell = min(shares_to_sell, current_volume)

        if self.benchmark:
            slippage, market_impact = self.benchmark.compute_components(alpha=0.001, shares=int(shares_to_sell))
            transaction_cost = slippage + market_impact
        else:
            transaction_cost = shares_to_sell * current_price * 0.001

        self.remaining_shares -= shares_to_sell
        self.current_step += 1

        done = self.current_step >= self.trading_horizon or self.remaining_shares <= 0
        reward = -transaction_cost

```

```

        return self._get_observation(), reward, done, {}

data = pd.read_csv('AAPL_Quotes_Data.csv')
benchmark = Benchmark(data)
env = DummyVecEnv([lambda: TradingEnv(data, benchmark=benchmark)])

sac_model = SAC('MlpPolicy', env, verbose=1)
sac_model.learn(total_timesteps=100000)
sac_model.save('sac_trading_model')
obs = env.reset()
timestamps_sac, share_sizes_sac, slippage_costs_sac, market_impacts_sac = [], [], [], []

for i in range(env.get_attr('trading_horizon')[0]):
    action, _ = sac_model.predict(obs, deterministic=True)
    obs, _, done, _ = env.step(action)
    timestamps_sac.append(data.iloc[i]['timestamp'])
    shares_sold = max(0.01, action[0]) * env.get_attr('remaining_shares')[0]
    share_sizes_sac.append(shares_sold)
    slippage, market_impact = benchmark.compute_components(alpha=0.001, shares=int(shares_sold), idx=i)
    slippage_costs_sac.append(slippage)
    market_impacts_sac.append(market_impact)
    if done:
        break

sac_trade_schedule = pd.DataFrame({'timestamp': timestamps_sac, 'shares_to_sell': share_sizes_sac, 'slippage': slippage_costs_sac, 'market_impact': market_impacts_sac})

ppo_model = PPO('MlpPolicy', env, verbose=1)
ppo_model.learn(total_timesteps=100000)
ppo_model.save('ppo_trading_model')

obs = env.reset()
timestamps_ppo, share_sizes_ppo, slippage_costs_ppo, market_impacts_ppo = [], [], [], []

for i in range(env.get_attr('trading_horizon')[0]):
    action, _ = ppo_model.predict(obs, deterministic=True)
    obs, _, done, _ = env.step(action)
    timestamps_ppo.append(data.iloc[i]['timestamp'])
    shares_sold = max(0.01, action[0]) * env.get_attr('remaining_shares')[0]
    share_sizes_ppo.append(shares_sold)
    slippage, market_impact = benchmark.compute_components(alpha=0.001, shares=int(shares_sold), idx=i)
    slippage_costs_ppo.append(slippage)
    market_impacts_ppo.append(market_impact)
    if done:
        break

ppo_trade_schedule = pd.DataFrame({'timestamp': timestamps_ppo, 'shares_to_sell': share_sizes_ppo, 'slippage': slippage_costs_ppo, 'market_impact': market_impacts_ppo})

twap_trades = benchmark.get_twap_trades(data, initial_inventory=1000)
vwap_trades = benchmark.get_vwap_trades(data, initial_inventory=1000)

twap_slippage, twap_market_impact = benchmark.simulate_strategy(twap_trades, data, preferred_timeframe='1m')
vwap_slippage, vwap_market_impact = benchmark.simulate_strategy(vwap_trades, data, preferred_timeframe='1m')

sac_total_cost = sum(slippage_costs_sac) + sum(market_impacts_sac)
ppo_total_cost = sum(slippage_costs_ppo) + sum(market_impacts_ppo)
twap_total_cost = sum(twap_slippage) + sum(twap_market_impact)
vwap_total_cost = sum(vwap_slippage) + sum(vwap_market_impact)

```

```
print(f'SAC Total Cost: {sac_total_cost}')  
print(f'PPO Total Cost: {ppo_total_cost}')  
print(f'TWAP Total Cost: {twap_total_cost}')  
print(f'VWAP Total Cost: {vwap_total_cost}')
```